

Godavari Foundation's

# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 1

## Subject: 440(B) Lab on Big Data Analytics

**Title:** Use Hadoop File System (HDFS) commands to perform basic operations like creating directories, uploading files, listing files, and deleting files in HDFS.

- **Hadoop Installation:** Ensure you have Hadoop installed and configured correctly.
- **Hadoop Daemons Running:** Start the HDFS daemons (NameNode and DataNode) using the `start-dfs.sh` script (usually found in the `sbin` directory of your Hadoop installation).

### HDFS Commands and Steps:

#### 1. Creating a Directory:

- Command: `hdfs dfs -mkdir <directory_path>`
- Explanation: This command creates a new directory in HDFS.
- Steps:
  - Open your terminal or command prompt.
  - Type: `hdfs dfs -mkdir /my_hdfs_directory` (Replace `/my_hdfs_directory` with the desired path).
  - Press Enter.

#### 2. Uploading a File from Local File System to HDFS:

- Command: `hdfs dfs -put <local_file_path> <hdfs_destination_path>`
- Explanation: This command copies a file from your local file system to HDFS.
- Steps:
  - Open your terminal.
  - Type: `hdfs dfs -put /path/to/local/file.txt /my_hdfs_directory/file.txt` (Replace `/path/to/local/file.txt` with the actual path to your local file and `/my_hdfs_directory/file.txt` with the desired HDFS destination).
  - Press Enter.

#### 3. Listing Files and Directories in HDFS:

- Command: `hdfs dfs -ls <hdfs_path>`
- Explanation: This command lists the files and directories present at the specified path in HDFS.
- Steps:
  - Open your terminal.
  - To list the contents of the root directory: `hdfs dfs -ls /`

- To list the contents of the directory created previously: `hdfs dfs -ls /my_hdfs_directory`
  - Press Enter.
4. **Deleting a File or Directory in HDFS:**
- Command: `hdfs dfs -rm <hdfs_file_path>` (for files)
  - Command: `hdfs dfs -rm -r <hdfs_directory_path>` (for directories, `-r` for recursive deletion)
  - Explanation: These commands remove files or directories from HDFS. Use `-r` to delete directories and their contents.
  - Steps:
    - Open your terminal.
    - To delete a file: `hdfs dfs -rm /my_hdfs_directory/file.txt`
    - To delete a directory and its content: `hdfs dfs -rm -r /my_hdfs_directory`
    - Press Enter.
5. **Viewing the content of a file:**
- Command: `hdfs dfs -cat <hdfs_file_path>`
  - Explanation: This command displays the contents of the specified file on the console.
  - Steps:
    - Open your terminal.
    - Type: `hdfs dfs -cat /my_hdfs_directory/file.txt`
    - Press Enter.

### Important Notes:

- **HDFS Paths:** HDFS paths always start with `/`.
- **Case Sensitivity:** HDFS paths are case-sensitive.
- **Error Handling:** If you encounter errors, check the Hadoop logs for more details.
- **Hadoop User:** Ensure you are running the commands as the Hadoop user or a user with sufficient permissions.
- **Web UI:** You can also use the Hadoop web UI (usually accessible at `http://<namenode_hostname>:9870`) to browse and manage HDFS files.

### Practical Steps:

1. **Open your terminal.**
2. **Create an HDFS directory:**

Bash

```
hdfs dfs -mkdir /my_practice_dir
```

- This creates a directory named `my_practice_dir` in the root of your HDFS.

3. **Upload a local file to HDFS:**

Bash

```
hdfs dfs -put ~/test.txt /my_practice_dir/uploaded_test.txt
```

- This copies the `test.txt` file from your home directory to the `my_practice_dir` directory in HDFS, renaming it to `uploaded_test.txt`.

**4. List the contents of the HDFS directory:**

Bash

```
hdfs dfs -ls /my_practice_dir
```

- You should see `uploaded_test.txt` listed.

**5. View the content of the uploaded file:**

Bash

```
hdfs dfs -cat /my_practice_dir/uploaded_test.txt
```

- The contents of your `test.txt` file will be displayed in the terminal.

**6. Delete the uploaded file:**

Bash

```
hdfs dfs -rm /my_practice_dir/uploaded_test.txt
```

- This will remove `uploaded_test.txt` from HDFS.

**7. Verify the file is deleted:**

Bash

```
hdfs dfs -ls /my_practice_dir
```

- The file should no longer be listed.

**8. Delete the HDFS directory:**

Bash

```
hdfs dfs -rm -r /my_practice_dir
```

- The `-r` option is crucial for deleting directories.

**9. Verify the directory is deleted:**

Bash

```
hdfs dfs -ls /
```

- The directory `my_practice_dir` should no longer be present in the root of HDFS.



# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 2

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Implement a Java program to interact with HDFS (reading and writing files).

## Steps to Implement Java HDFS Interaction:

1. **Setup Hadoop:** Ensure Hadoop is running.
2. **Create Java Project:** In your IDE or with a build tool.
3. **Add Hadoop Dependency:** Include `hadoop-client` in your project's dependencies (e.g., `pom.xml` for Maven).
4. **Write Java Code:** Copy and paste the provided Java code into your project.
5. **Configure Hadoop (if needed):** Set `fs.defaultFS` in the `Configuration` object if not automatically detected.
6. **Adjust File Paths:** Modify `hdfsWritePath` and `hdfsReadPath` in the code.
7. **Compile Java Code.**
8. **Run Java Code.**
9. **Verify in HDFS:** Use `hdfs dfs -ls` and `hdfs dfs -cat` to check the file.

## Code:

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FSDataInputStream;
import org.apache.hadoop.fs.FSDataOutputStream;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import java.io.BufferedReader;
import java.io.IOException;
import java.io.InputStreamReader;

public class HdfsInteraction {
    public static void main(String[] args) {
        try {
            Configuration conf = new Configuration();
            // Optional: If needed, set your HDFS configuration here
            // conf.set("fs.defaultFS", "hdfs://localhost:9000");
            FileSystem fs = FileSystem.get(conf);
            // Write to HDFS
            Path writePath = new Path("/my_java_hdfs_file.txt");
            try (FSDataOutputStream out = fs.create(writePath)) {
                out.writeBytes("Hello, HDFS from Java!\n");
            }
        } catch (IOException e) {
            e.printStackTrace();
        }
    }
}
```

```
        out.writeBytes("Another line written via Java.\n");
        System.out.println("Written to: " + writePath);
    }
    // Read from HDFS
    Path readPath = new Path("/my_java_hdfs_file.txt");
    if (fs.exists(readPath)) {
        try (FSDataInputStream in = fs.open(readPath);
            BufferedReader reader = new BufferedReader(new InputStreamReader(in))) {
            String line;
            System.out.println("Reading from: " + readPath);
            while ((line = reader.readLine()) != null) {
                System.out.println(line);
            }
        }
    } else {
        System.out.println("File not found: " + readPath);
    }
    fs.close();
} catch (IOException e) {
    e.printStackTrace();
}
}
```

# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 3

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Use Hadoop's built-in commands to manage files and directories.

## **Hadoop File/Directory Management Steps (using `hadoop fs`):**

1. **Open Terminal:** Access your command-line interface.
2. **Listing Files/Directories:**
  - Type: `hadoop fs -ls <path>` (replace <path>)
  - Press Enter.
3. **Creating Directories:**
  - Type: `hadoop fs -mkdir -p <path>` (replace <path>)
  - Press Enter.
4. **Uploading Files:**
  - Type: `hadoop fs -put <local_path> <hdfs_path>` (replace paths)
  - Press Enter.
5. **Downloading Files:**
  - Type: `hadoop fs -get <hdfs_path> <local_path>` (replace paths)
  - Press Enter.
6. **Moving Files:**
  - Type: `hadoop fs -mv <source_hdfs_path> <destination_hdfs_path>` (replace paths)
  - Press Enter.
7. **Copying Files:**
  - Type: `hadoop fs -cp <source_hdfs_path> <destination_hdfs_path>` (replace paths)
  - Press Enter.
8. **Removing Files/Directories:**
  - Type: `hadoop fs -rm -r <hdfs_path>` (replace <path>)
  - Press Enter. (Use with caution!)
9. **Displaying File Content:**
  - Type: `hadoop fs -cat <hdfs_path>` (replace <path>)
  - Press Enter.

# Listing files/directories

```
hadoop fs -ls /user/yourusername/
```

# Creating directories

```
hadoop fs -mkdir -p /user/yourusername/mydir/subdir
```

# Uploading files

```
hadoop fs -put localfile.txt /user/yourusername/mydir/
```

# Downloading files

```
hadoop fs -get /user/yourusername/mydir/localfile.txt downloadedfile.txt
```

# Moving files

```
hadoop fs -mv /user/yourusername/mydir/localfile.txt /user/yourusername/newdir/
```

# Copying files

```
hadoop fs -cp /user/yourusername/mydir/localfile.txt /user/yourusername/copieddir/
```

# Removing files/directories

```
hadoop fs -rm -r /user/yourusername/mydir/
```

# Displaying file content

```
hadoop fs -cat /user/yourusername/mydir/localfile.txt
```



# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 4

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Implement Map Side Join and Reduce Side Join.

## Assumptions:

- You have a Hadoop cluster set up.
- You have two datasets that you want to join.
- You are familiar with basic MapReduce concepts.

## Dataset Example:

Let's say we have two datasets:

- **users.txt:** Contains user IDs and names.
- **transactions.txt:** Contains transaction IDs, user IDs, and amounts.

## 1. Map-Side Join:

### When to Use:

- When one of the datasets is small enough to fit into the memory of each mapper.
- When the datasets are pre-partitioned and sorted by the join key.

### Steps:

#### 1. Prepare the Small Dataset:

- Ensure the smaller dataset (`users.txt` in our example) is available in HDFS.
- If needed, pre-process the small dataset so it can be easily loaded into a data structure (e.g., a `HashMap`).

#### 2. Implement the Mapper:

- In the `setup()` method of your mapper, load the small dataset into a data structure (e.g., a `HashMap`).
- In the `map()` method:
  - Read a record from the larger dataset (`transactions.txt`).
  - Extract the join key (user ID).
  - Look up the corresponding record in the loaded small dataset.
  - Join the records and emit the result.

#### 3. Configure the Job:

- Use the Distributed Cache to distribute the small dataset to all mappers.

- Ensure that the input data for the larger dataset is partitioned and sorted by the join key, in the same way the small dataset is sorted.
  - Set the number of reducers to 0, because this is a map only job.
4. **Run the Job:**
- Execute the Hadoop job.

## 2. Reduce-Side Join:

### When to Use:

- When both datasets are large.
- When the datasets are not pre-partitioned or sorted.

### Steps:

1. **Implement the Mappers:**
  - Create two mappers, one for each dataset.
  - In each mapper:
    - Read a record.
    - Extract the join key.
    - Emit the join key and the record, along with a tag indicating the dataset it came from.
2. **Implement the Reducer:**
  - The reducer receives all records with the same join key.
  - Separate the records based on their tags.
  - Perform the join operation.
  - Emit the joined records.
3. **Configure the Job:**
  - Set the input paths for both datasets.
  - Set the mapper classes.
  - Set the reducer class.
  - Set the partitioner and grouping comparator so that all keys are sent to the same reducer, and that values are grouped correctly.
4. **Run the Job:**
  - Execute the Hadoop job.

### Code:

```
// Map-Side Join Code (Conceptual)

import java.io.BufferedReader;
import java.io.FileReader;
import java.io.IOException;
import java.net.URI;
import java.util.HashMap;
import java.util.Map;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.filecache.DistributedCache;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
```

```

import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MapSideJoin {
    public static class MapSideJoinMapper extends Mapper<LongWritable, Text, Text, Text> {
        private Map<String, String> users = new HashMap<>();

        @Override
        protected void setup(Context context) throws IOException, InterruptedException {
            Path[] cacheFiles = DistributedCache.getLocalCacheFiles(context.getConfiguration());
            if (cacheFiles != null && cacheFiles.length > 0) {
                String line;
                BufferedReader reader = new BufferedReader(new
                FileReader(cacheFiles[0].toString()));
                while ((line = reader.readLine()) != null) {
                    String[] parts = line.split(",");
                    users.put(parts[0], parts[1]); // userID, userName
                }
                reader.close();
            }
        }

        @Override
        protected void map(LongWritable key, Text value, Context context) throws IOException,
        InterruptedException {
            String[] parts = value.toString().split(","); // transactionID, userID, amount
            String userID = parts[1];
            String userName = users.get(userID);
            if (userName != null) {
                context.write(new Text(userID), new Text(parts[0] + "," + userName + "," +
                parts[2]));
            }
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "MapSideJoin");
        job.setJarByClass(MapSideJoin.class);
        job.setMapperClass(MapSideJoinMapper.class);
        job.setNumReduceTasks(0); // Map-only job
        FileInputFormat.addInputPath(job, new Path(args[0])); // transactions.txt
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        DistributedCache.addCacheFile(new URI(args[2]), job.getConfiguration()); // users.txt
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

// Reduce-Side Join Code (Conceptual)
import java.io.IOException;
import java.util.ArrayList;
import java.util.List;

```

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class ReduceSideJoin {

    public static class UserMapper extends Mapper<LongWritable, Text, Text, Text> {
        @Override
        protected void map(LongWritable key, Text value, Context context) throws IOException,
            InterruptedException {
            String[] parts = value.toString().split(",");
            context.write(new Text(parts[0]), new Text("U," + parts[1])); // userID, "U,userName"
        }
    }

    public static class TransactionMapper extends Mapper<LongWritable, Text, Text, Text> {
        @Override
        protected void map(LongWritable key, Text value, Context context) throws IOException,
            InterruptedException {
            String[] parts = value.toString().split(",");
            context.write(new Text(parts[1]), new Text("T," + parts[0] + "," + parts[2])); // userID,
            "T,transactionID,amount"
        }
    }

    public static class ReduceSideJoinReducer extends Reducer<Text, Text, Text, Text> {
        @Override
        protected void reduce(Text key, Iterable<Text> values, Context context) throws
            IOException, InterruptedException {
            List<String> users = new ArrayList<>();
            List<String> transactions = new ArrayList<>();
            for (Text value : values) {
                String[] parts = value.toString().split(",");
                if (parts[0].equals("U")) {
                    users.add(parts[1]);
                } else {
                    transactions.add(parts[1] + "," + parts[2]);
                }
            }
            for (String user : users) {
                for (String transaction : transactions) {
                    context.write(key, new Text(user + "," + transaction));
                }
            }
        }
    }
}

```

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "ReduceSideJoin");
    job.setJarByClass(ReduceSideJoin.class);
    job.setMapperClass(UserMapper.class);
    job.setMapperClass(TransactionMapper.class);
    job.setReducerClass(ReduceSideJoinReducer.class);
    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(Text.class);

    FileInputFormat.addInputPath(job, new Path(args[0])); // users.txt
    FileInputFormat.addInputPath(job, new Path(args[1])); // transactions.txt
    FileOutputFormat.setOutputPath(job, new Path(args[2]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```



# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 5

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Implement Secondary Sorting.

- **Define Composite Key:**

- Create a Java class (e.g., `CompositeKey.java`) that implements `WritableComparable`.
- Add fields for your natural key and secondary key.
- Implement `compareTo()`: Compare natural keys first, then secondary keys if natural keys are equal.
- Implement `readFields()` and `write()` for serialization.

- **Implement Mapper:**

- Create your mapper class.
- In the `map()` method:
  - Extract the natural key and secondary key from your input data.
  - Create an instance of your `CompositeKey` class.
  - Emit the `CompositeKey` as the key and your value.

- **Create Partitioner (Recommended):**

- Create a custom partitioner class (e.g., `NaturalKeyPartitioner.java`) that extends `Partitioner`.
- In the `getPartition()` method:
  - Extract only the natural key from the `CompositeKey`.
  - Use the natural key to determine the partition.

- **Create Grouping Comparator:**

- Create a custom grouping comparator class (e.g., `NaturalKeyGroupingComparator.java`) that extends `WritableComparator`.
- In the `compare()` method:
  - Extract only the natural key from the `CompositeKey` instances.
  - Compare only the natural keys.

- **Implement Reducer:**

- Create your reducer class.

- In the `reduce()` method:
  - The values associated with each natural key will be sorted by the secondary key.
  - Process the sorted values.
- **Configure Hadoop Job:**
  - Create a `Job` configuration.
  - Set the mapper class, reducer class, and input/output formats.
  - Set the `MapOutputKeyClass` to your `CompositeKey` class.
  - Set the partitioner class:
 

```
job.setPartitionerClass(NaturalKeyPartitioner.class);
```
  - Set the grouping comparator class:
 

```
job.setGroupingComparatorClass(NaturalKeyGroupingComparator.class);
```
  - Set the jar by class, and input and output paths.
- **Compile and Package:**
  - Compile your Java classes using `javac`.
  - Create a JAR file containing your compiled classes.
- **Run the Job:**
  - Use the `hadoop jar` command to run your MapReduce job on your Hadoop cluster.
  - Provide the input and output paths as arguments.

```
import java.io.DataInput;
import java.io.DataOutput;
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.io.WritableComparator;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Partitioner;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class SecondarySort {

    // Composite Key Class
    public static class CompositeKey implements WritableComparable<CompositeKey> {
        private Text naturalKey;
        private IntWritable secondaryKey;
```



```

public CompositeKey() {
    this.naturalKey = new Text();
    this.secondaryKey = new IntWritable();
}

public CompositeKey(Text naturalKey, IntWritable secondaryKey) {
    this.naturalKey = naturalKey;
    this.secondaryKey = secondaryKey;
}

@Override
public void write(DataOutput out) throws IOException {
    naturalKey.write(out);
    secondaryKey.write(out);
}

@Override
public void readFields(DataInput in) throws IOException {
    naturalKey.readFields(in);
    secondaryKey.readFields(in);
}

@Override
public int compareTo(CompositeKey other) {
    int naturalKeyComparison = this.naturalKey.compareTo(other.naturalKey);
    if (naturalKeyComparison == 0) {
        return this.secondaryKey.compareTo(other.secondaryKey);
    }
    return naturalKeyComparison;
}

public Text getNaturalKey() {
    return naturalKey;
}

public IntWritable getSecondaryKey() {
    return secondaryKey;
}
}

// Mapper Class
public static class SecondarySortMapper extends Mapper<LongWritable, Text,
CompositeKey, IntWritable> {
    private CompositeKey compositeKey = new CompositeKey();
    private IntWritable valueWritable = new IntWritable();

    @Override
    protected void map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException {
        String[] parts = value.toString().split(",");
        Text naturalKey = new Text(parts[0]);
        int secondaryKey = Integer.parseInt(parts[1]);
    }
}

```

```

        int valueInt = Integer.parseInt(parts[2]);

        compositeKey.naturalKey.set(naturalKey);
        compositeKey.secondaryKey.set(secondaryKey);
        valueWritable.set(valueInt);

        context.write(compositeKey, valueWritable);
    }
}

// Partitioner Class
public static class NaturalKeyPartitioner extends Partitioner<CompositeKey, IntWritable>
{
    @Override
    public int getPartition(CompositeKey key, IntWritable value, int numPartitions) {
        return (key.getNaturalKey().hashCode() & Integer.MAX_VALUE) % numPartitions;
    }
}

// Grouping Comparator Class
public static class NaturalKeyGroupingComparator extends WritableComparator {
    protected NaturalKeyGroupingComparator() {
        super(CompositeKey.class, true);
    }

    @Override
    public int compare(WritableComparable w1, WritableComparable w2) {
        CompositeKey key1 = (CompositeKey) w1;
        CompositeKey key2 = (CompositeKey) w2;
        return key1.getNaturalKey().compareTo(key2.getNaturalKey());
    }
}

// Reducer Class
public static class SecondarySortReducer extends Reducer<CompositeKey, IntWritable,
Text, Text> {
    @Override
    protected void reduce(CompositeKey key, Iterable<IntWritable> values, Context
context) throws IOException, InterruptedException {
        StringBuilder valueString = new StringBuilder();
        for (IntWritable value : values) {
            valueString.append(value.get()).append(",");
        }
        context.write(key.getNaturalKey(), new Text(valueString.toString()));
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "SecondarySort");
    job.setJarByClass(SecondarySort.class);
    job.setMapperClass(SecondarySortMapper.class);

```

```
job.setReducerClass(SecondarySortReducer.class);
job.setPartitionerClass(NaturalKeyPartitioner.class);
job.setGroupingComparatorClass(NaturalKeyGroupingComparator.class);
job.setMapOutputKeyClass(CompositeKey.class);
job.setMapOutputValueClass(IntWritable.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(Text.class);

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```



# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 6

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Pipeline multiple Map Reduce jobs.

- **Develop Individual MapReduce Jobs:**

- Create each MapReduce job as a separate Java class.
- Ensure each job reads input from and writes output to HDFS.
- Test each job individually to ensure it works correctly.

- **Determine Job Dependency and Order:**

- Identify the sequence in which the jobs must execute.
- Understand the data flow: how the output of one job serves as the input to the next.

- **Implement a Driver Class (or Script):**

- Create a driver class (or a shell script) to manage the execution of the jobs.
- This driver will:
  - Configure and run the first MapReduce job.
  - Wait for the first job to complete.
  - Configure and run the second MapReduce job, using the output of the first job as its input.
  - Repeat for any subsequent jobs.

- **Handle Intermediate Output:**

- Choose a location in HDFS to store the intermediate output of each job.
- Make sure the next job in the pipeline reads from this location.
- Consider deleting intermediate output if it's no longer needed to save space.

- **Error Handling:**

- Implement error checking in your driver class or script.
- Check the return code of each job.
- If a job fails, decide whether to stop the pipeline or take other actions.

- **Job Configuration and Parameterization:**

- If necessary, parameterize your jobs to allow for flexible configuration.
- Use command-line arguments or configuration files to pass parameters to the jobs.

- This is especially helpful when dealing with different input/output paths.
- **Testing:**
  - Thoroughly test the entire pipeline.
  - Start with small datasets and gradually increase the size.
  - Verify the output of each job in the pipeline.

## Code:

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MapReducePipeline {

    // Job 1: Word Count
    public static class Job1Mapper extends Mapper<LongWritable, Text, Text,
IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, Context context) throws
java.io.IOException, InterruptedException {
            String line = value.toString();
            String[] words = line.split("\\s+");
            for (String w : words) {
                word.set(w);
                context.write(word, one);
            }
        }
    }

    public static class Job1Reducer extends Reducer<Text, IntWritable, Text, IntWritable>
    {
        public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
java.io.IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }
}
```

```

// Job 2: Filter words with count > 2
public static class Job2Mapper extends Mapper<LongWritable, Text, Text,
IntWritable> {
    public void map(LongWritable key, Text value, Context context) throws
java.io.IOException, InterruptedException {
        String[] parts = value.toString().split("\\t");
        Text word = new Text(parts[0]);
        IntWritable count = new IntWritable(Integer.parseInt(parts[1]));
        context.write(word, count);
    }
}

public static class Job2Reducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
java.io.IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        if (sum > 2) {
            context.write(key, new IntWritable(sum));
        }
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    // Job 1 Configuration and Execution
    Job job1 = Job.getInstance(conf, "Word Count");
    job1.setJarByClass(MapReducePipeline.class);
    job1.setMapperClass(Job1Mapper.class);
    job1.setReducerClass(Job1Reducer.class);
    job1.setOutputKeyClass(Text.class);
    job1.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job1, new Path(args[0])); // Input text file
    FileOutputFormat.setOutputPath(job1, new Path("/intermediate/wordcount")); //
Intermediate output

    if (!job1.waitForCompletion(true)) {
        System.err.println("Job 1 failed!");
        System.exit(1);
    }

    // Job 2 Configuration and Execution
    Job job2 = Job.getInstance(conf, "Filter Words");
    job2.setJarByClass(MapReducePipeline.class);
    job2.setMapperClass(Job2Mapper.class);
    job2.setReducerClass(Job2Reducer.class);
    job2.setOutputKeyClass(Text.class);

```

```
        job2.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job2, new Path("/intermediate/wordcount")); // Input
from Job 1 output
        FileOutputFormat.setOutputPath(job2, new Path(args[1])); // Final output

        if (!job2.waitForCompletion(true)) {
            System.err.println("Job 2 failed!");
            System.exit(1);
        }

        System.exit(0);
    }
}
```



# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 7

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Create and use UDFs in Pig Latin scripts.

**1. Write the UDF in Java:**

- Create a Java class that extends `org.apache.pig.EvalFunc<ReturnType>`.
- Implement the `exec(Tuple input)` method.
- This method takes a tuple as input and returns a value of the specified `ReturnType`.
- Handle potential exceptions (e.g., `IOException`).
- Compile the Java class into a JAR file.

**2. Register the JAR in Pig:**

- In your Pig script, use the `REGISTER` command to register the JAR file containing your UDF.
- This makes the UDF available for use in the script.

**3. Use the UDF in Your Pig Script:**

- Call the UDF in your Pig script like any other built-in function.
- Pass the required arguments to the UDF.

**Example Scenario:**

Let's create a UDF that converts a string to uppercase.

**1. Create the Java UDF (UppercaseUDF.java):**

Java

```
import java.io.IOException;
import org.apache.pig.EvalFunc;
import org.apache.pig.data.Tuple;

public class UppercaseUDF extends EvalFunc<String> {

    public String exec(Tuple input) throws IOException {
        if (input == null || input.size() == 0) {
            return null;
        }
        try {
            String inputString = (String) input.get(0);
            if (inputString == null) {
                return null;
            }
        }
    }
}
```

```

        return inputString.toUpperCase();
    } catch (Exception e) {
        throw new IOException("Caught exception processing input row ",
e);
    }
}
}

```

## 2. Compile the Java UDF:

Bash

```

javac -classpath $(pig classpath) UppercaseUDF.java
jar cf uppercaseudf.jar UppercaseUDF.class

```

## 3. Use the UDF in a Pig Script (uppercase.pig):

Code snippet

```

-- Load data
data = LOAD 'input.txt' AS (text:chararray);

-- Register the UDF JAR
REGISTER 'uppercaseudf.jar';

-- Apply the UDF
uppercase_data = FOREACH data GENERATE UppercaseUDF(text);

-- Store the result
STORE uppercase_data INTO 'output';

```

## Example Input Data (input.txt):

```

hello world
pig latin
hadoop ecosystem

```

## 4. Run the Pig Script:

Bash

```

pig uppercase.pig

```

## Explanation:

- The UppercaseUDF class extends EvalFunc<String>, indicating that it returns a string.
- The exec() method takes a tuple as input, extracts the string, and returns its uppercase version.
- The Pig script registers the uppercaseudf.jar file.
- The FOREACH statement applies the UppercaseUDF to the text field of each record.
- The result is stored in the output directory.

## Important Notes:

- **Error Handling:** Always include proper error handling in your UDFs to prevent unexpected failures.
- **Data Types:** Ensure that the data types in your UDF match the data types in your Pig script.

- **Complex UDFs:** For more complex UDFs, you can use other Pig interfaces like `Algebraic` and `Accumulator` to implement aggregate functions and incremental computations.
- **Testing:** Thoroughly test your UDFs before using them in production.
- **Classpath:** When compiling your UDF, make sure to include the Pig classpath to access the necessary Pig classes.
- **Data Structures:** Pig Tuples, Bags, and Maps can be accessed within the UDF.
- **Return types:** Ensure your return type matches the <Return Type> of the `EvalFunc`.

## Code:

```
// UppercaseUDF.java
import java.io.IOException;
import org.apache.pig.EvalFunc;
import org.apache.pig.data.Tuple;

public class UppercaseUDF extends EvalFunc<String> {

    public String exec(Tuple input) throws IOException {
        if (input == null || input.size() == 0) {
            return null;
        }
        try {
            String inputString = (String) input.get(0);
            if (inputString == null) {
                return null;
            }
            return inputString.toUpperCase();
        } catch (Exception e) {
            throw new IOException("Caught exception processing input row ", e);
        }
    }
}

-- uppercase.pig
-- Load data
data = LOAD 'input.txt' AS (text:chararray);

-- Register the UDF JAR
REGISTER 'uppercaseudf.jar';

-- Apply the UDF
uppercase_data = FOREACH data GENERATE UppercaseUDF(text);

-- Store the result
STORE uppercase_data INTO 'output';
```



# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 8

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Integrate UDFs to enhance the functionality of Pig scripts.

- **Identify Functionality to Enhance:**

- Determine specific tasks in your Pig script that could benefit from custom logic.
- Common scenarios include data validation, complex transformations, or external API calls.

- **Develop the UDF in Java:**

- Create a Java class that extends `org.apache.pig.EvalFunc<ReturnType>`.
- Implement the `exec(Tuple input)` method to perform the desired task.
- Handle potential exceptions and edge cases.
- Compile the UDF into a JAR file.

- **Register the UDF JAR in Pig:**

- Add the `REGISTER 'path/to/your/udf.jar';` command to your Pig script.
- Replace `'path/to/your/udf.jar'` with the actual path to your JAR file.

- **Incorporate the UDF into Your Pig Script:**

- Call the UDF within your Pig script using the UDF's class name as a function.
- Pass the required arguments to the UDF.
- Use the UDF's output in subsequent Pig operations.

```
// EmailValidation.java
import java.io.IOException;
import java.util.regex.Matcher;
import java.util.regex.Pattern;
import org.apache.pig.EvalFunc;
import org.apache.pig.data.Tuple;
```

```
public class EmailValidation extends EvalFunc<String> {

    private static final Pattern EMAIL_PATTERN = Pattern.compile(
        "^[A-Za-z0-9+_.-]+@[A-Za-z0-9.-]+$");

    public String exec(Tuple input) throws IOException {
```

```

    if (input == null || input.size() == 0) {
        return null;
    }
    try {
        String email = (String) input.get(0);
        if (email == null) {
            return null;
        }
        Matcher matcher = EMAIL_PATTERN.matcher(email);
        if (matcher.matches()) {
            return email.toLowerCase();
        } else {
            return "INVALID";
        }
    } catch (Exception e) {
        throw new IOException("Error validating email", e);
    }
}
}

```

-- validate\_emails.pig

-- Load data

emails = LOAD 'emails.txt' AS (email:chararray);

-- Register the UDF JAR

REGISTER 'emailvalidation.jar';

-- Validate and transform emails

validated\_emails = FOREACH emails GENERATE EmailValidation(email);

-- Filter out invalid emails

valid\_emails = FILTER validated\_emails BY \$0 != 'INVALID';

-- Store the results

STORE valid\_emails INTO 'valid\_email\_output';

# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 9

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Implement and execute HiveQL queries to perform data retrieval and manipulation.

Let's implement and execute HiveQL queries to perform data retrieval and manipulation.

## Steps:

### 1. Start the Hive CLI:

- Open your terminal.
- Type hive and press Enter. This will launch the Hive command-line interface.

### 2. Create a Database (Optional):

- If you want to organize your tables, create a database:

SQL

```
CREATE DATABASE mydatabase;
```

```
USE mydatabase;
```

### 3. Create Tables:

- Define the structure of your data by creating tables.
- Specify the data types and delimiters.
- Example:

SQL

```
CREATE TABLE users (
```

```
  id INT,
```

```
  name STRING,
```

```
  age INT,
```

```
  city STRING
```

```
)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ','
```

```
STORED AS TEXTFILE;
```

```
CREATE TABLE sales (
```

```
  sale_id INT,
```

```
  user_id INT,
```

```
  product STRING,
```

```
  amount DOUBLE
```

```
)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ','
```

```
STORED AS TEXTFILE;
```

#### 4. Load Data into Tables:

- Use the LOAD DATA LOCAL INPATH command to load data from your local file system into the Hive tables.
- Example:

SQL

```
LOAD DATA LOCAL INPATH '/path/to/users.csv' OVERWRITE INTO TABLE users;  
LOAD DATA LOCAL INPATH '/path/to/sales.csv' OVERWRITE INTO TABLE sales;
```

- Replace /path/to/users.csv and /path/to/sales.csv with the actual paths to your CSV files.

#### 5. Perform Data Retrieval (SELECT Queries):

- Use SELECT statements to retrieve data from the tables.
- Examples:

- Retrieve all data from the users table:

SQL

```
SELECT * FROM users;
```

- Retrieve specific columns:

SQL

```
SELECT name, city FROM users;
```

- Apply filters using WHERE clauses:

SQL

```
SELECT * FROM users WHERE age > 30;
```

- Aggregate data using GROUP BY and aggregate functions:

SQL

```
SELECT city, AVG(age) FROM users GROUP BY city;
```

- Join tables:

SQL

```
SELECT u.name, s.product, s.amount  
FROM users u  
JOIN sales s ON u.id = s.user_id;
```

- Order the results:

SQL

```
SELECT * from sales ORDER BY amount DESC;
```

- Use LIMIT to limit the number of rows:



## SQL

```
SELECT * FROM users LIMIT 10;
```

### 6. Perform Data Manipulation (INSERT/UPDATE/DELETE):

#### ○ INSERT:

- Insert data from one table into another:

## SQL

```
INSERT INTO TABLE new_users SELECT * FROM users WHERE age > 25;
```

- Insert single rows: (Hive 0.14.0 and later)

## SQL

```
INSERT INTO users VALUES (6,'John Doe',45,'London');
```

#### ○ UPDATE/DELETE:

- Hive does not directly support UPDATE and DELETE in the same way as traditional relational databases.
- You can achieve similar results by creating new tables with the desired modifications or by using partitions.
- For example, to "delete" rows, you can create a new table with only the rows you want to keep:

## SQL

```
CREATE TABLE users_filtered AS SELECT * FROM users WHERE age <= 30;
```

- Then you can drop the original table and rename the new table.

### 7. Exit the Hive CLI:

- Type quit; or exit; and press Enter.

### Example CSV Data (users.csv):

#### Code snippet

```
1,Alice,25,New York
2,Bob,32,London
3,Charlie,40,Paris
4,David,28,Tokyo
5,Eve,35,Sydney
```

### Example CSV Data (sales.csv):

#### Code snippet

```
101,1,Laptop,1200.00
102,2,Phone,800.00
103,1,Tablet,500.00
104,3,Desktop,1500.00
105,4,Monitor,300.00
```

```
-- Start Hive CLI: hive

-- Create a database (optional)
CREATE DATABASE mydatabase;
USE mydatabase;

-- Create tables
CREATE TABLE users (
    id INT,
    name STRING,
    age INT,
    city STRING
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;

CREATE TABLE sales (
    sale_id INT,
    user_id INT,
    product STRING,
    amount DOUBLE
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;

-- Load data (replace paths)
LOAD DATA LOCAL INPATH '/path/to/users.csv' OVERWRITE INTO TABLE users;
LOAD DATA LOCAL INPATH '/path/to/sales.csv' OVERWRITE INTO TABLE sales;

-- Data retrieval (SELECT queries)
SELECT * FROM users;
SELECT name, city FROM users;
SELECT * FROM users WHERE age > 30;
SELECT city, AVG(age) FROM users GROUP BY city;
SELECT u.name, s.product, s.amount FROM users u JOIN sales s ON u.id = s.user_id;
SELECT * FROM sales ORDER BY amount DESC;
SELECT * FROM users LIMIT 10;

-- Data manipulation (INSERT/UPDATE/DELETE)
INSERT INTO TABLE new_users SELECT * FROM users WHERE age > 25;

-- (Hive 0.14.0 and later)
INSERT INTO users VALUES (6,'John Doe',45,'London');

-- "DELETE" (create new table)
CREATE TABLE users_filtered AS SELECT * FROM users WHERE age <= 30;

-- Exit Hive CLI
quit; -- or exit;
```

# Godavari Institute Of Management & Research, Jalgaon

Name: \_\_\_\_\_

Roll No: \_\_\_\_\_

Date of Performance: \_\_/\_\_/20\_\_

Batch: \_\_\_\_\_

Class: M.C.A. (I) Practical no: 10

**Subject: 440(B) Lab on Big Data Analytics**

**Title:** Perform operations like joins, group by, and aggregations in Hive.

Let's perform joins, group by, and aggregations in Hive.

**Steps:**

**1. Start Hive CLI:**

- Open your terminal.
- Type `hive` and press Enter.

**2. Create/Use a Database (Optional):**

- If you want to organize your tables, create a database:

**SQL**

```
CREATE DATABASE mydatabase;  
USE mydatabase;
```

**3. Create Tables:**

- Create tables with data you want to use.
- Example:

**SQL**

```
CREATE TABLE users (  
    id INT,  
    name STRING,  
    city STRING  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;  
  
CREATE TABLE orders (  
    order_id INT,  
    user_id INT,  
    product STRING,  
    quantity INT,  
    price DOUBLE  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','
```

STORED AS TEXTFILE;

#### 4. Load Data:

- Load data into the tables from local files.
- Example:

##### SQL

```
LOAD DATA LOCAL INPATH '/path/to/users.csv' OVERWRITE INTO  
TABLE users;  
LOAD DATA LOCAL INPATH '/path/to/orders.csv' OVERWRITE INTO  
TABLE orders;
```

- Replace /path/to/users.csv and /path/to/orders.csv with the actual paths.

#### 5. Joins:

- **Inner Join:**

##### SQL

```
SELECT u.name, o.product  
FROM users u  
JOIN orders o ON u.id = o.user_id;
```

- **Left Outer Join:**

##### SQL

```
SELECT u.name, o.product  
FROM users u  
LEFT OUTER JOIN orders o ON u.id = o.user_id;
```

- **Right Outer Join:**

##### SQL

```
SELECT u.name, o.product  
FROM users u  
RIGHT OUTER JOIN orders o ON u.id = o.user_id;
```

- **Full Outer Join:**

##### SQL

```
SELECT u.name, o.product  
FROM users u  
FULL OUTER JOIN orders o ON u.id = o.user_id;
```

#### 6. Group By:

- Group data by a column.
- Example:

##### SQL

```
SELECT city, COUNT(*)
```

```
FROM users
GROUP BY city;
```

## 7. Aggregations:

- Use aggregate functions with GROUP BY.
- Examples:

- **Count:**

SQL

```
SELECT product, COUNT(*)
FROM orders
GROUP BY product;
```

- **Sum:**

SQL

```
SELECT user_id, SUM(price * quantity)
FROM orders
GROUP BY user_id;
```

- **Average:**

SQL

```
SELECT product, AVG(price)
FROM orders
GROUP BY product;
```

- **Max:**

SQL

```
SELECT user_id, MAX(price * quantity)
FROM orders
GROUP BY user_id;
```

- **Min:**

SQL

```
SELECT user_id, MIN(price * quantity)
FROM orders
GROUP BY user_id;
```

- **Combining Aggregations:**

SQL

```
SELECT city, COUNT(*), AVG(id)
FROM users
GROUP BY city;
```

## 8. Filtering with HAVING:

- Use `HAVING` to filter aggregated results.
- Example:

## SQL

```
SELECT product, COUNT(*)
FROM orders
GROUP BY product
HAVING COUNT(*) > 1;
```

### 9. Exit Hive CLI:

- Type `quit;` or `exit;` and press Enter.

### Example CSV Data (users.csv):

#### Code snippet

```
1,Alice,New York
2,Bob,London
3,Charlie,Paris
4,David,New York
```

### Example CSV Data (orders.csv):

#### Code snippet

```
101,1,Laptop,1,1200.00
102,2,Phone,2,800.00
103,1,Tablet,1,500.00
104,3,Desktop,1,1500.00
105,4,Monitor,2,300.00
106,1,Mouse,3,25.00
```

-- Start Hive CLI: `hive`

-- Create/Use a database (optional)

```
CREATE DATABASE mydatabase;
USE mydatabase;
```

-- Create tables

```
CREATE TABLE users (
  id INT,
  name STRING,
  city STRING
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;
```

```
CREATE TABLE orders (
```

```
  order_id INT,
  user_id INT,
  product STRING,
```

```
    quantity INT,  
    price DOUBLE  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;  
  
-- Load data (replace paths)  
LOAD DATA LOCAL INPATH '/path/to/users.csv' OVERWRITE INTO TABLE users;  
LOAD DATA LOCAL INPATH '/path/to/orders.csv' OVERWRITE INTO TABLE orders;  
  
-- Joins  
SELECT u.name, o.product FROM users u JOIN orders o ON u.id = o.user_id;  
SELECT u.name, o.product FROM users u LEFT OUTER JOIN orders o ON u.id =  
o.user_id;  
SELECT u.name, o.product FROM users u RIGHT OUTER JOIN orders o ON u.id =  
o.user_id;  
SELECT u.name, o.product FROM users u FULL OUTER JOIN orders o ON u.id =  
o.user_id;  
  
-- Group By  
SELECT city, COUNT(*) FROM users GROUP BY city;  
  
-- Aggregations  
SELECT product, COUNT(*) FROM orders GROUP BY product;  
SELECT user_id, SUM(price * quantity) FROM orders GROUP BY user_id;  
SELECT product, AVG(price) FROM orders GROUP BY product;  
SELECT user_id, MAX(price * quantity) FROM orders GROUP BY user_id;  
SELECT user_id, MIN(price * quantity) FROM orders GROUP BY user_id;  
SELECT city, COUNT(*), AVG(id) FROM users GROUP BY city;  
  
-- Filtering with HAVING  
SELECT product, COUNT(*) FROM orders GROUP BY product HAVING COUNT(*) > 1;  
  
-- Exit Hive CLI  
quit; -- or exit;
```